

# BTC price analysis

```
In [55]:  
#!/usr/bin/env python  
#pip install yfinance  
import yfinance as yf  
import datetime as dt  
import pandas as pd  
  
start = dt.datetime(2010,3,1)  
end = dt.datetime(2024,3,2)  
  
df = yf.download('BTC-AUD', start, end)
```

```
[*****100%*****] 1 of 1 completed
```

```
In [56]: df.tail()
```

```
Out[56]:
```

| Date       | Open         | High         | Low          | Close        | Adj Close    | V      |
|------------|--------------|--------------|--------------|--------------|--------------|--------|
| 2024-02-26 | 78665.468750 | 79180.101562 | 78279.570312 | 78820.445312 | 78820.445312 | 23483  |
| 2024-02-27 | 78816.343750 | 84033.539062 | 77795.875000 | 83375.875000 | 83375.875000 | 52106  |
| 2024-02-28 | 83371.226562 | 87965.406250 | 83398.804688 | 87236.039062 | 87236.039062 | 76036  |
| 2024-02-29 | 87214.226562 | 98355.726562 | 86821.304688 | 96242.187500 | 96242.187500 | 128168 |
| 2024-03-01 | 96233.554688 | 97557.101562 | 93069.312500 | 94151.382812 | 94151.382812 | 100764 |

```
In [57]: # df.describe()  
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
DatetimeIndex: 3454 entries, 2014-09-17 to 2024-03-01  
Data columns (total 6 columns):  
 #   Column      Non-Null Count  Dtype     
---  --          --          --  
 0   Open        3454 non-null    float64  
 1   High        3454 non-null    float64  
 2   Low         3454 non-null    float64  
 3   Close       3454 non-null    float64  
 4   Adj Close   3454 non-null    float64  
 5   Volume      3454 non-null    int64  
dtypes: float64(5), int64(1)  
memory usage: 188.9 KB
```

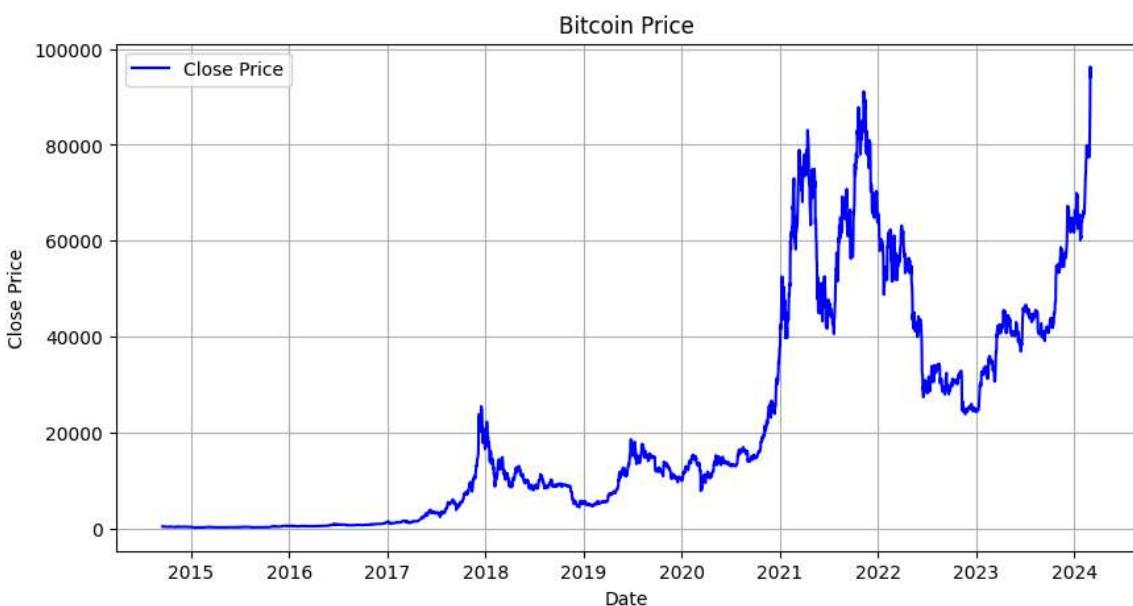
## Simple BTC graphing

## BTC PRICE

In [58]:

```
import matplotlib.pyplot as plt

plt.figure(figsize=(10, 5))
plt.plot(df.index, df['Close'], linestyle='-', color='b', label='Close Price')
plt.title('Bitcoin Price')
plt.xlabel('Date')
plt.ylabel('Close Price')
plt.legend()
plt.grid(True)
plt.show()
```



## BTC logged price

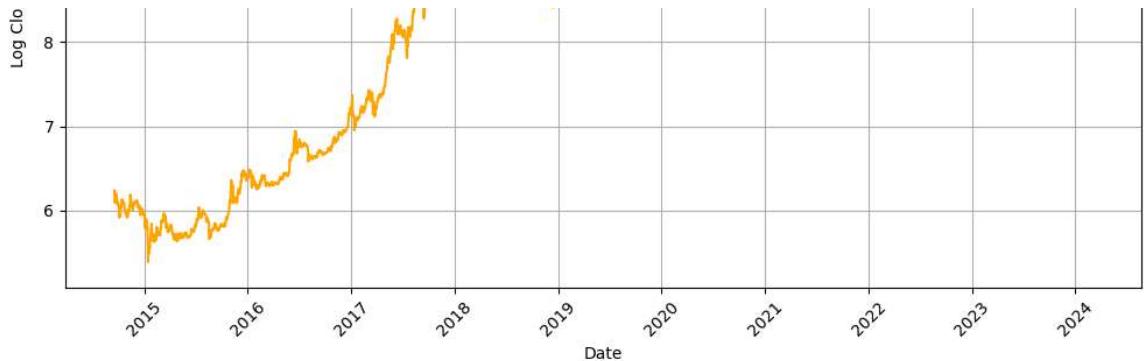
In [59]:

```
import numpy as np

df['Log_Close'] = np.log(df['Close'])

# Plotting the logarithm of the close price
plt.figure(figsize=(10, 6))
plt.plot(df.index, df['Log_Close'], label='Log Close Price', color='orange')
plt.title('Logarithm of Close Price Over Time')
plt.xlabel('Date')
plt.ylabel('Log Close Price')
plt.legend()
plt.grid(True)
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```





## Checking for autocorrelation

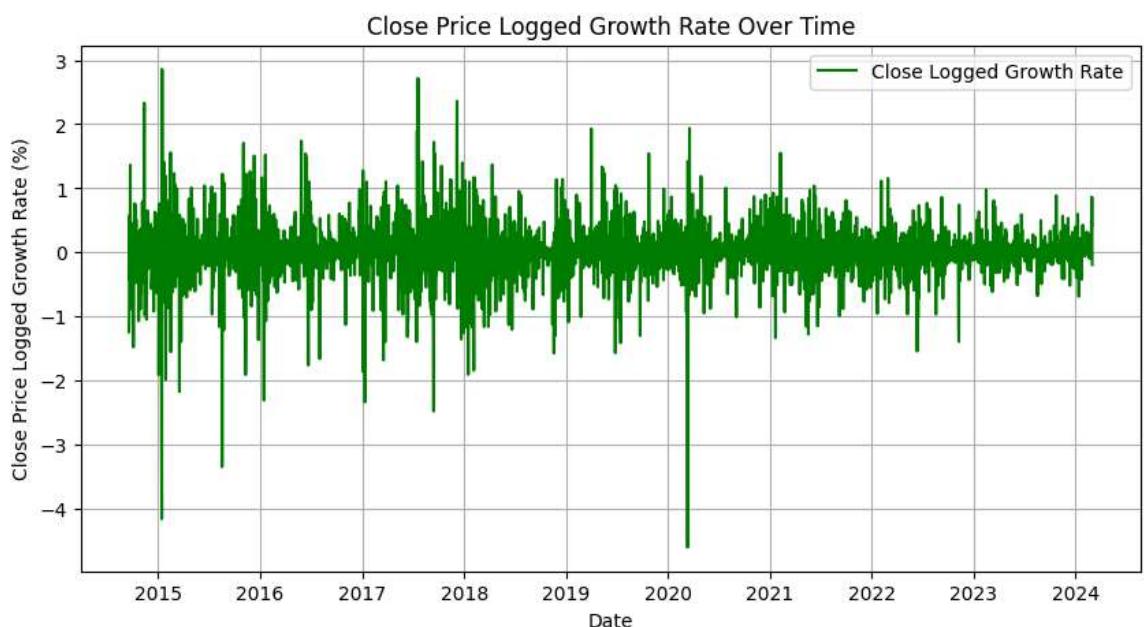
### Making data stationary for unbiased estimates

We take the I(1) (integrated to the first order) approach by differencing (i.e. taking the growth rate of logged prices).

In [60]:

```
# Calculate the daily percentage change (growth rate)
df['Close_Log_Growth_Rate'] = df['Log_Close'].pct_change() * 100
# Also make just the growth rate for later analysis
df['Close_Growth_Rate'] = df['Close'].pct_change() * 100

# Plotting growth rate on its own
plt.figure(figsize=(10, 5))
plt.plot(df.index, df['Close_Log_Growth_Rate'], linestyle='-', color='g', lab
plt.title('Close Price Logged Growth Rate Over Time')
plt.xlabel('Date')
plt.ylabel('Close Price Logged Growth Rate (%)')
plt.legend()
plt.grid(True)
plt.show()
```



In [61]:

```
import pandas as pd
from statsmodels.tsa.stattools import adfuller
```

```

# Assuming df is your DataFrame and 'Close_Log_Growth_Rate' is the column you
# Perform Augmented Dickey-Fuller test
result = adfuller(df['Close_Log_Growth_Rate'].dropna())

# Output the results
print('ADF Statistic: %f' % result[0])
print('p-value: %f' % result[1])
print('Critical Values:')
for key, value in result[4].items():
    print('\t%s: %.3f' % (key, value))

# Interpretation
if result[1] > 0.05:
    print("The series is not stationary")
else:
    print("The series is stationary")

```

ADF Statistic: -18.433370  
p-value: 0.000000  
Critical Values:  
1%: -3.432  
5%: -2.862  
10%: -2.567  
The series is stationary

## Deciding upon sample

We decide to try 2 different samples to test for any autocorrelation in BTC price (as it is difficult to test daily prices over a decade of time):

1. Daily data over 1 year: for short term fluctuation analysis
2. Monthly data for 5 years: for long term trend analysis
3. Intraday data over 1 week (recorded at 5 minute intervals): for extremely short-term analysis

In [69]:

```

#!/usr/bin/env python3
import yfinance as yf
import datetime as dt
import pandas as pd

# 1. Daily data over 1 year
daily_data_1yr = df.last('365D')

# 2. Monthly data for 5 years
monthly_data_5yr = df.last('5Y').resample('M').mean()

# 3. Intra-day data for 1 week
start_forIntraDay = dt.datetime(2024,2,24)
end_forIntraDay = dt.datetime(2024,3,2)
minute_data_1wk = yf.download('BTC-AUD', start_forIntraDay, end_forIntraDay,
# Calculate Logged price
minute_data_1wk['Log_Close'] = np.log(minute_data_1wk['Close'])
# Calculate the daily percentage change (growth rate)
minute_data_1wk['Close_Log_Growth_Rate'] = minute_data_1wk['Log_Close'].pct_change()
# Also make just the growth rate alone
minute_data_1wk['Close_Growth_Rate'] = minute_data_1wk['Close'].pct_change()

```

```
# Drop all rows with NaN values
minute_data_1wk.dropna(inplace=True)

daily_data_1yr.head(), monthly_data_5yr.head(), minute_data_1wk.head()
```

```
/var/folders/1c/xdsm7htj4334pn2n9vx5kz90000gn/T/ipykernel_1757/2681600539.py:8: FutureWarning: last is deprecated and will be removed in a future version. Please create a mask and filter using `loc` instead
    daily_data_1yr = df.last('365D')
/var/folders/1c/xdsm7htj4334pn2n9vx5kz90000gn/T/ipykernel_1757/2681600539.py:12: FutureWarning: last is deprecated and will be removed in a future version. Please create a mask and filter using `loc` instead
    monthly_data_5yr = df.last('5Y').resample('M').mean()
[*****100%*****] 1 of 1 completed
```

```
Out[69]: (
```

|            | Open         | High         | Low          | Close        | \ |
|------------|--------------|--------------|--------------|--------------|---|
| Date       |              |              |              |              |   |
| 2023-03-03 | 34875.359375 | 34879.886719 | 32809.906250 | 33041.796875 |   |
| 2023-03-04 | 33042.156250 | 33102.437500 | 32797.796875 | 33025.867188 |   |
| 2023-03-05 | 33027.039062 | 33410.500000 | 32957.601562 | 33217.597656 |   |
| 2023-03-06 | 33219.523438 | 33526.722656 | 33073.078125 | 33317.765625 |   |
| 2023-03-07 | 33315.636719 | 33416.160156 | 33403.851562 | 33722.808594 |   |

|            | Adj Close    | Volume      | Log_Close | Close_Log_Growth_Rate | \ |
|------------|--------------|-------------|-----------|-----------------------|---|
| Date       |              |             |           |                       |   |
| 2023-03-03 | 33041.796875 | 38508297433 | 10.405529 | -0.515872             |   |
| 2023-03-04 | 33025.867188 | 16497181114 | 10.405046 | -0.004634             |   |
| 2023-03-05 | 33217.597656 | 19716899742 | 10.410835 | 0.055633              |   |
| 2023-03-06 | 33317.765625 | 25776901263 | 10.413846 | 0.028922              |   |
| 2023-03-07 | 33722.808594 | 34550990217 | 10.425930 | 0.116035              |   |

|            | Close_Growth_Rate |
|------------|-------------------|
| Date       |                   |
| 2023-03-03 | -5.252771         |
| 2023-03-04 | -0.048211         |
| 2023-03-05 | 0.580546          |
| 2023-03-06 | 0.301551          |
| 2023-03-07 | 1.215697 ,        |

|            | Open         | High         | Low          | Close        | \ |
|------------|--------------|--------------|--------------|--------------|---|
| Date       |              |              |              |              |   |
| 2020-01-31 | 12107.797915 | 12369.311460 | 11944.036920 | 12228.686114 |   |
| 2020-02-29 | 14475.996633 | 14676.984813 | 14208.783574 | 14450.521518 |   |
| 2020-03-31 | 11088.148374 | 11463.077589 | 10589.622102 | 10995.810342 |   |
| 2020-04-30 | 11363.521647 | 11663.056380 | 11206.143978 | 11458.122363 |   |
| 2020-05-31 | 14190.745432 | 14515.087922 | 13897.235383 | 14217.975334 |   |

|            | Adj Close    | Volume       | Log_Close | Close_Log_Growth_Rate | \ |
|------------|--------------|--------------|-----------|-----------------------|---|
| Date       |              |              |           |                       |   |
| 2020-01-31 | 12228.686114 | 4.010258e+10 | 9.407679  | 0.106693              |   |
| 2020-02-29 | 14450.521518 | 6.023657e+10 | 9.577553  | -0.019830             |   |
| 2020-03-31 | 10995.810342 | 6.723576e+10 | 9.292311  | -0.073974             |   |
| 2020-04-30 | 11458.122363 | 6.113336e+10 | 9.344805  | 0.085897              |   |
| 2020-05-31 | 14217.975334 | 6.379741e+10 | 9.561507  | 0.021925              |   |

|            | Close_Growth_Rate |
|------------|-------------------|
| Date       |                   |
| 2020-01-31 | 1.041995          |
| 2020-02-29 | -0.161160         |
| 2020-03-31 | -0.319569         |
| 2020-04-30 | 0.854873          |
| 2020-05-31 | 0.256802 ,        |

|                           | Open         | High         | Low          | \ |
|---------------------------|--------------|--------------|--------------|---|
| Datetime                  |              |              |              |   |
| 2024-02-24 00:02:00+00:00 | 77425.343750 | 77425.343750 | 77425.343750 |   |

```

2024-02-24 00:04:00+00:00 77469.875000 77469.875000 77450.648438
2024-02-24 00:05:00+00:00 77422.781250 77422.781250 77422.781250
2024-02-24 00:06:00+00:00 77422.781250 77422.781250 77406.968750
2024-02-24 00:07:00+00:00 77390.695312 77390.695312 77390.695312

          Close   Adj Close   Volume Log_Close
\
Datetime
2024-02-24 00:02:00+00:00 77425.343750 77425.343750      0 11.257069
2024-02-24 00:04:00+00:00 77450.648438 77450.648438 8751104 11.257396
2024-02-24 00:05:00+00:00 77422.781250 77422.781250 14110720 11.257036
2024-02-24 00:06:00+00:00 77406.968750 77406.968750 487424 11.256832
2024-02-24 00:07:00+00:00 77390.695312 77390.695312      0 11.256622

          Close_Log_Growth_Rate  Close_Growth_Rate
Datetime
2024-02-24 00:02:00+00:00           0.003017  0.033966
2024-02-24 00:04:00+00:00           0.002903  0.032683
2024-02-24 00:05:00+00:00          -0.003197 -0.035981
2024-02-24 00:06:00+00:00          -0.001814 -0.020424
2024-02-24 00:07:00+00:00          -0.001868 -0.021023 )

```

## Testing for autocorrelation and modelling

### Short term (daily data over 1 year)

In [63]:

```

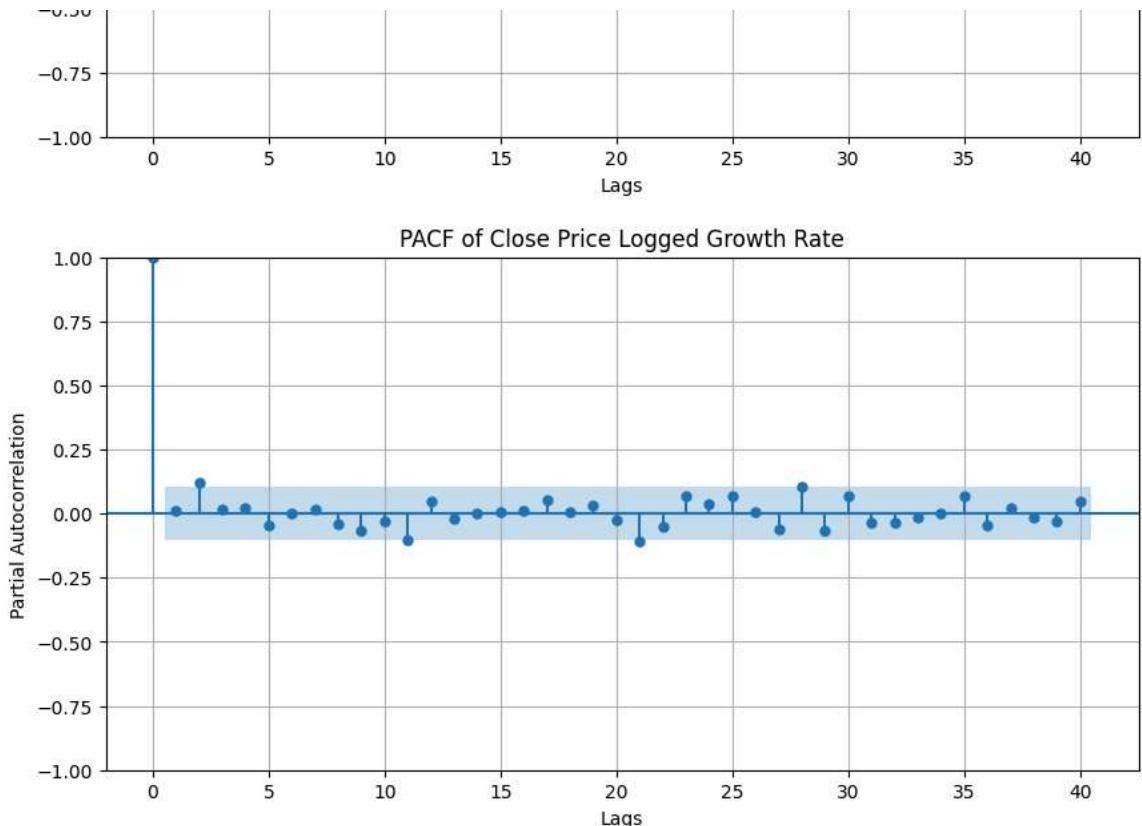
import pandas as pd
import matplotlib.pyplot as plt
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf

# ACF plot
plt.figure(figsize=(10, 5))
plot_acf(daily_data_1yr['Close_Log_Growth_Rate'].dropna(), ax=plt.gca(), lags=20)
plt.title('ACF of Close Price Logged Growth Rate')
plt.xlabel('Lags')
plt.ylabel('Autocorrelation')
plt.grid(True)
plt.show()

# PACF plot
plt.figure(figsize=(10, 5))
plot_pacf(daily_data_1yr['Close_Log_Growth_Rate'].dropna(), ax=plt.gca(), lags=20)
plt.title('PACF of Close Price Logged Growth Rate')
plt.xlabel('Lags')
plt.ylabel('Partial Autocorrelation')
plt.grid(True)
plt.show()

```





The fact that both ACF and PACF truncate after the second lag might suggest that a mixed ARMA model is not necessary, and a simpler AR(p) or MA(q) model might be adequate for modeling the time series.

As the second lag is significant in the ACF and PACF, we test for autocorrelation using an AR(2) and MA(2) and test for which has a lower AIC to measure model suitability.

### Testing AR(2)

```
In [64]: import pandas as pd
from statsmodels.tsa.arima.model import ARIMA

# Initialise and fit the ARIMA model with order (2,0,0) for AR(2) process
model = ARIMA(daily_data_1yr['Close_Log_Growth_Rate'].dropna(), order=(2,0,0))
model_fit = model.fit()

# Print out the summary of the model
print(model_fit.summary())

# You can also make predictions
# For example, to forecast the next 5 values:
forecast = model_fit.forecast(steps=5)
print(forecast)
```

```
/Library/Frameworks/Python.framework/Versions/3.11/lib/python3.11/site-packages/statsmodels/tsa/base/tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency D will be used.
    self._init_dates(dates, freq)
/Library/Frameworks/Python.framework/Versions/3.11/lib/python3.11/site-packages/statsmodels/tsa/base/tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency D will be used.
    self._init_dates(dates, freq)
/Library/Frameworks/Python.framework/Versions/3.11/lib/python3.11/site-packages/statsmodels/tsa/base/tsa_model.py:473: ValueWarning: No frequency information
```

```

was provided, so inferred frequency D will be used.
self._init_dates(dates, freq)
                                         SARIMAX Results
=====
Dep. Variable:      Close_Log_Growth_Rate   No. Observations:                  3
65
Model:                  ARIMA(2, 0, 0)    Log Likelihood               48.7
24
Date:          Sun, 03 Mar 2024     AIC                   -89.4
48
Time:              16:07:52        BIC                   -73.8
48
Sample:         03-03-2023     HQIC                  -83.2
48
                           - 03-01-2024
Covariance Type:            opg
=====
            coef      std err       z     P>|z|      [0.025      0.975]
-----
const      0.0251      0.013     1.902     0.057     -0.001      0.051
ar.L1      0.0095      0.047     0.201     0.840     -0.083      0.102
ar.L2      0.1275      0.055     2.338     0.019      0.021      0.234
sigma2     0.0448      0.002    18.888     0.000      0.040      0.049
=====
=====
Ljung-Box (L1) (Q):           0.00  Jarque-Bera (JB):                 10
6.84
Prob(Q):                      0.97  Prob(JB):
0.00
Heteroskedasticity (H):       0.94  Skew:
0.41
Prob(H) (two-sided):          0.73  Kurtosis:
5.52
=====
=====

Warnings:
[1] Covariance matrix calculated using the outer product of gradients (complex-step).
2024-03-02    0.129966
2024-03-03   -0.001538
2024-03-04    0.038201
2024-03-05    0.021807
2024-03-06    0.026720
Freq: D, Name: predicted_mean, dtype: float64

```

The first lag is insignificant, but the second lag is significant. To exclude the first lag, we'd have to include the second lag as basically an exogenous variable. This is difficult to interpret in a model, so I will refrain from doing so.

## Testing MA(2)

In [65]:

```

import pandas as pd
from statsmodels.tsa.arima.model import ARIMA

# Initialise and fit the ARIMA model with order (2,0,0) for AR(2) process
model = ARIMA(daily_data_1yr['Close_Log_Growth_Rate'].dropna(), order=(0,0,2))
model_fit = model.fit()

# Print out the summary of the model
print(model_fit.summary())

```

```

# You can also make predictions
# For example, to forecast the next 5 values:
forecast = model_fit.forecast(steps=5)
print(forecast)



/Library/Frameworks/Python.framework/Versions/3.11/lib/python3.11/site-packages/statsmodels/tsa/base/tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency D will be used.
    self._init_dates(dates, freq)
/Library/Frameworks/Python.framework/Versions/3.11/lib/python3.11/site-packages/statsmodels/tsa/base/tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency D will be used.
    self._init_dates(dates, freq)
/Library/Frameworks/Python.framework/Versions/3.11/lib/python3.11/site-packages/statsmodels/tsa/base/tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency D will be used.
    self._init_dates(dates, freq)
                                          SARIMAX Results
=====
==

Dep. Variable:      Close_Log_Growth_Rate   No. Observations:            3
65
Model:                  ARIMA(0, 0, 2)    Log Likelihood             48.5
48
Date:          Sun, 03 Mar 2024     AIC                  -89.0
97
Time:                  16:07:53       BIC                  -73.4
97
Sample:          03-03-2023     HQIC                 -82.8
97
                           - 03-01-2024
Covariance Type:            opg
=====
                                         coef      std err      z      P>|z|      [0.025      0.975]
-----
const        0.0251        0.013     1.963      0.050      3.95e-05      0.050
ma.L1        0.0041        0.047     0.089      0.929      -0.087      0.095
ma.L2        0.1201        0.053     2.245      0.025      0.015      0.225
sigma2       0.0449        0.002    19.087      0.000      0.040      0.049
=====
==

Ljung-Box (L1) (Q):           0.00  Jarque-Bera (JB):            11
0.32
Prob(Q):                   0.95  Prob(JB):
0.00
Heteroskedasticity (H):      0.94  Skew:
0.42
Prob(H) (two-sided):         0.71  Kurtosis:
5.56
=====
==

Warnings:
[1] Covariance matrix calculated using the outer product of gradients (complex-step).
2024-03-02    0.117883
2024-03-03   -0.006695
2024-03-04    0.025051
2024-03-05    0.025051
2024-03-06    0.025051
Freq: D, Name: predicted_mean, dtype: float64

```

Similar results as with AR(2) Modelling on ARMA(2,2) doesn't help either

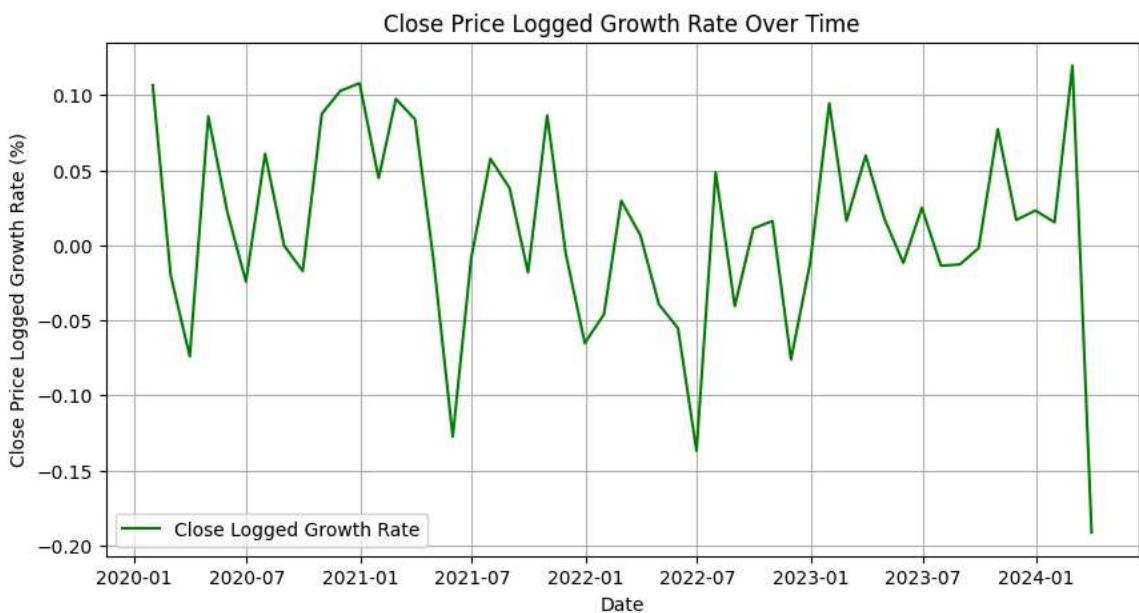
Similar results as with AR(2). Including an ARIMA(2,2) doesn't help either.

In sum, AR and MA modelling is not appropriate for residuals of BTC growth rate.

## Long term (monthly data over 5 years)

In [66]:

```
# Plotting growth rate on its own
plt.figure(figsize=(10, 5))
plt.plot(monthly_data_5yr.index, monthly_data_5yr['Close_Log_Growth_Rate'], 1
plt.title('Close Price Logged Growth Rate Over Time')
plt.xlabel('Date')
plt.ylabel('Close Price Logged Growth Rate (%)')
plt.legend()
plt.grid(True)
plt.show()
```



In [67]:

```
## Test for stationarity in this new df first:

import pandas as pd
from statsmodels.tsa.stattools import adfuller

# Perform Augmented Dickey-Fuller test
result = adfuller(monthly_data_5yr['Close_Log_Growth_Rate'].dropna())

# Output the results
print('ADF Statistic: %f' % result[0])
print('p-value: %f' % result[1])
print('Critical Values:')
for key, value in result[4].items():
    print('\t%s: %.3f' % (key, value))

# Interpretation
if result[1] > 0.05:
    print("The series is not stationary")
else:
    print("The series is stationary")
```

ADF Statistic: -6.075511

p-value: 0.000000

Critical Values:

1%: -3.568

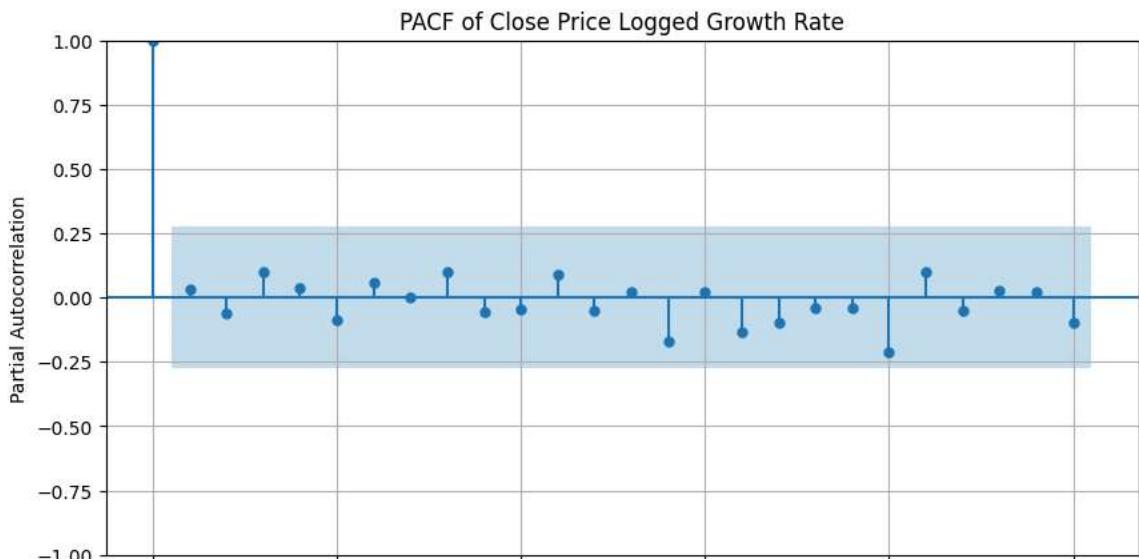
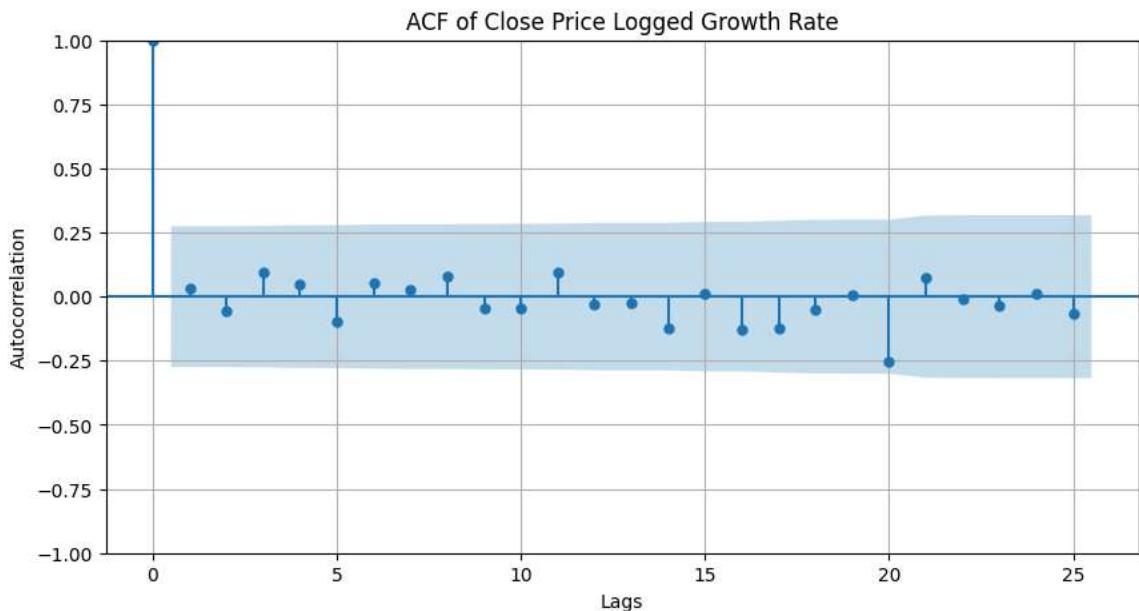
5%: -2.921  
10%: -2.599  
The series is stationary

In [68]:

```
import pandas as pd
import matplotlib.pyplot as plt
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf

# ACF plot
plt.figure(figsize=(10, 5))
plot_acf(monthly_data_5yr['Close_Log_Growth_Rate'].dropna(), ax=plt.gca(), l
plt.title('ACF of Close Price Logged Growth Rate')
plt.xlabel('Lags')
plt.ylabel('Autocorrelation')
plt.grid(True)
plt.show()

# PACF plot
plt.figure(figsize=(10, 5))
plot_pacf(monthly_data_5yr['Close_Log_Growth_Rate'].dropna(), ax=plt.gca(), l
plt.title('PACF of Close Price Logged Growth Rate')
plt.xlabel('Lags')
plt.ylabel('Partial Autocorrelation')
plt.grid(True)
plt.show()
```



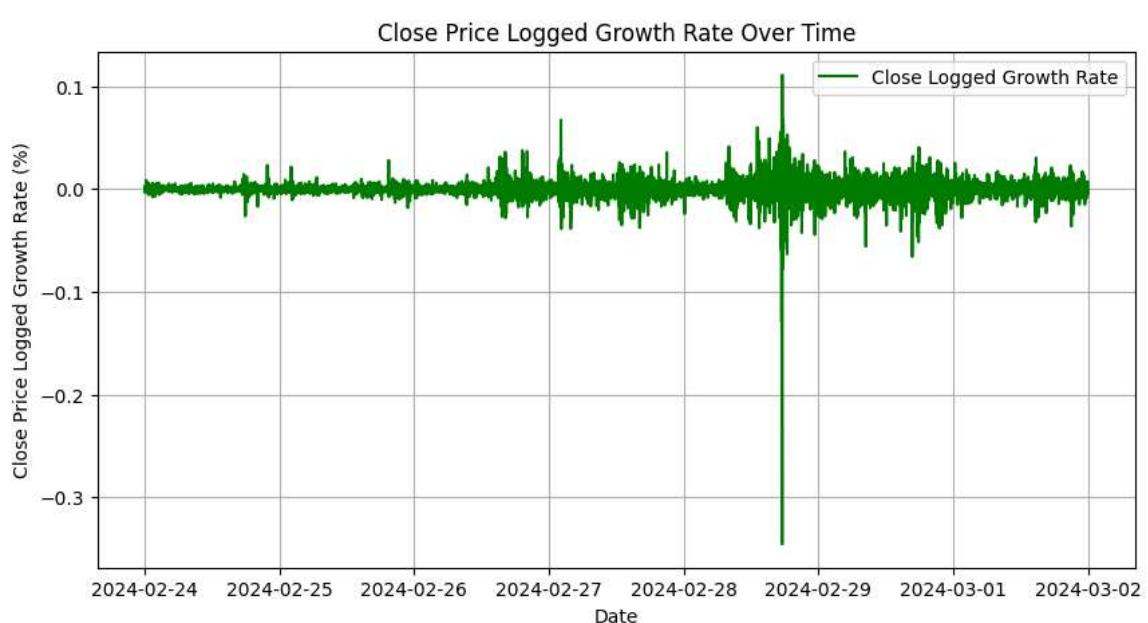


There is no autocorrelation at any of the lags. No autocorrelation modelling of the error term is required.

### Extreme short term (intra-day minute-by-minute)

In [70]:

```
# Plotting growth rate on its own
plt.figure(figsize=(10, 5))
plt.plot(minute_data_1wk.index, minute_data_1wk['Close_Log_Growth_Rate'], lin
plt.title('Close Price Logged Growth Rate Over Time')
plt.xlabel('Date')
plt.ylabel('Close Price Logged Growth Rate (%)')
plt.legend()
plt.grid(True)
plt.show()
```



In [71]:

```
## Test for stationarity in this new df first:

import pandas as pd
from statsmodels.tsa.stattools import adfuller

# Perform Augmented Dickey-Fuller test
result = adfuller(minute_data_1wk['Close_Log_Growth_Rate'].dropna())

# Output the results
print('ADF Statistic: %f' % result[0])
print('p-value: %f' % result[1])
print('Critical Values:')
for key, value in result[4].items():
    print('\t%s: %.3f' % (key, value))

# Interpretation
if result[1] > 0.05:
    print("The series is not stationary")
else:
    print("The series is stationary")
```

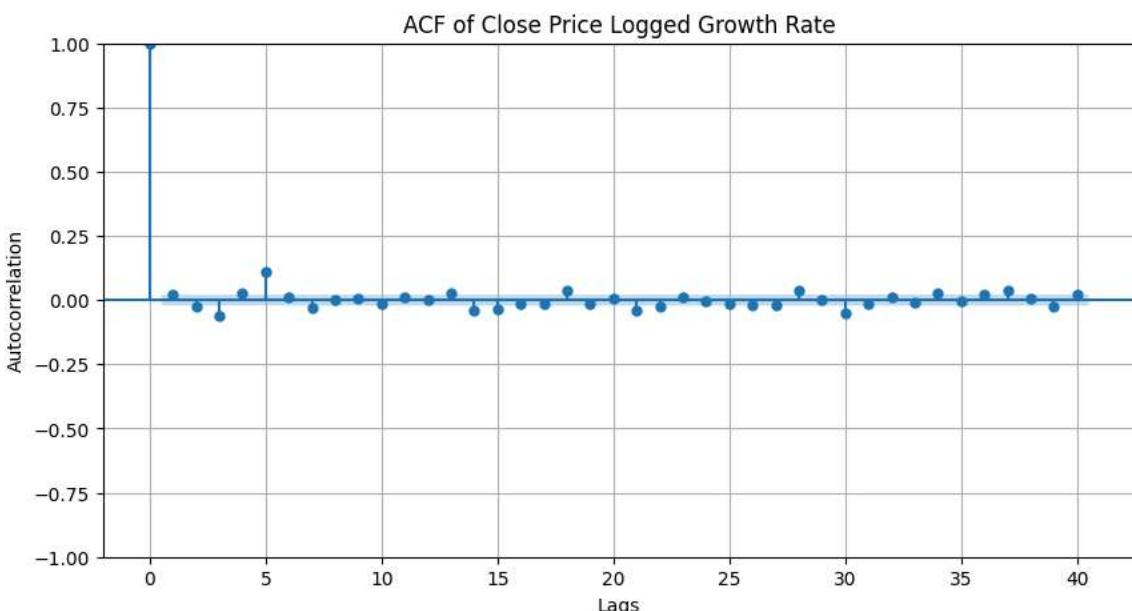
ADF Statistic: -15.502151

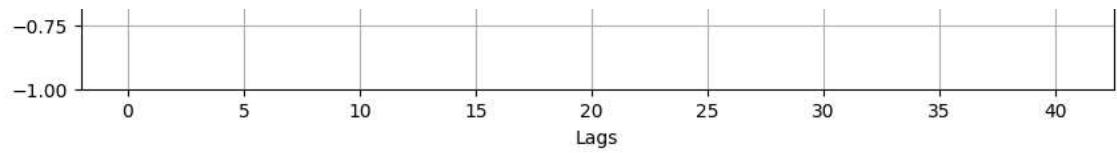
```
p-value: 0.000000
Critical Values:
    1%: -3.431
    5%: -2.862
    10%: -2.567
The series is stationary
```

```
In [73]: import pandas as pd
import matplotlib.pyplot as plt
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf

# ACF plot
plt.figure(figsize=(10, 5))
plot_acf(minute_data_1wk['Close_Log_Growth_Rate'].dropna(), ax=plt.gca(), lag=40)
plt.title('ACF of Close Price Logged Growth Rate')
plt.xlabel('Lags')
plt.ylabel('Autocorrelation')
plt.grid(True)
plt.show()

# PACF plot
plt.figure(figsize=(10, 5))
plot_pacf(minute_data_1wk['Close_Log_Growth_Rate'].dropna(), ax=plt.gca(), lag=40)
plt.title('PACF of Close Price Logged Growth Rate')
plt.xlabel('Lags')
plt.ylabel('Partial Autocorrelation')
plt.grid(True)
plt.show()
```





Both ACF and PACF are truncated and both have 5 lags that are significant. Both graphs exhibit a sine wave pattern.

## What actually affects the price of BTC if not its past values?